

Iyed Sebai

AI Engineer | RAG & Agentic Systems | Production LLM Platforms in Regulated Environments

iyed.sebai@gmail.com | +216 93 681 987 | github.com/The-iyed | linkedin.com/in/sebai-iyed

Tunis, Tunisia | Open to relocation to Saudi Arabia (sponsorship required) | Remote (GMT+1, full GST overlap)

Languages: Arabic (native) | English (fluent) | French (fluent)

Summary

AI engineer with 5+ years of software experience, focused on retrieval-augmented and agentic systems that must hold up under audit. I took a governed LLM platform from raw document corpora to production inside a Saudi government legal department, where it runs today across three services used daily by the department's legal staff. Depth in on-prem ingestion and PII-aware processing, hybrid retrieval with citation-level provenance, tool-calling agents with human-in-the-loop gates, and step-level tracing that lets compliance teams reconstruct any answer. Backed by a distributed-systems foundation (Go, Kafka, PostgreSQL) for the parts of AI work that are really systems work: throughput, reliability, and cost.

Technical Skills

AI / LLM	RAG architecture, hybrid search (vector + lexical), chunking & metadata strategy, grounded generation with citations, tool-calling agents, human-in-the-loop workflows, prompt engineering, evaluation & regression testing, hallucination guardrails
AI Tooling	OpenAI and Anthropic models, open-source/self-hosted models, LangChain, text-embedding-3-large, semantic reranking
Vector & Search	Azure AI Search (hybrid vector + keyword, semantic reranker), Neo4j (graph-backed retrieval)
Observability	Step-level agent tracing, usage analytics, retrieval coverage analysis, evaluation harnesses
Languages	Python, Go, TypeScript/JavaScript, Java
Backend	FastAPI, NestJS, Express, Gin; REST; WebSockets
Data	PostgreSQL, MongoDB, Neo4j; indexing & query optimisation
Infra	Azure, Docker, CI/CD; on-prem & air-gapped deployment

Experience

AI Engineer & Tech Lead — MOMRAH (Saudi Ministry) <i>Legal Department Remote via Tryloop</i>	Feb 2025 – Present <i>1 yr 7 mos</i>
---	---

- **Silver Stevie® Award**, 2026 Middle East & North Africa Stevie Awards — Excellence in Innovation in AI Products & Services, Government (for Sanad AI).
- Led design and delivery of **Sanad AI**, a governed LLM platform for a Saudi government legal department under strict security, compliance, and audit constraints; three services now run in daily production use.
- Shipped a **pleadings copilot** with human-in-the-loop review: evidence-grounded drafting, structured citations, and guardrails constraining generation to retrieved evidence to reduce hallucination.
- Built **tool-calling agents** querying internal legal sources to produce benchmark-style comparisons, each claim traceable to its supporting document.
- Delivered a **rulings analysis engine** generating legal KPIs and feeding learning signals back into the drafting workflow.
- Diagnosed a **retrieval defect** in which an over-aggressive character cap truncated judicial reasoning across the entire corpus, so the model never received a complete legal rationale; re-scoped extraction

and surfaced each ruling's statutory basis for the first time.

- Instrumented **step-level observability** across agent runs — prompts, retrieved evidence, tool calls, outputs — so legal and audit teams can reconstruct exactly how any response was produced.
- Contributed to the **Unified Platform** for the Minister and Executive Office: a role-aware decision-support interface synthesising KPIs, project status, regulatory dossiers, and correspondence with full source traceability.

Stack: Azure, Python, Docker, PostgreSQL/MongoDB, vector + lexical retrieval, LLM orchestration, evaluation & observability

AI Software Engineer — Builtop

Mar 2024 – Jan 2025

Remote / Saudi Arabia

11 mos

- Built a customer-facing **buyer assistant** for buy-now-pay-later users, answering eligibility, instalment-plan, and order-status questions grounded in live platform data.
- Built an **internal operations assistant** giving the ops team natural-language access to platform statistics, removing the ops team's dependency on manual reporting and ad-hoc SQL requests.
- Integrated both assistants into existing production systems and release processes within a 7-engineer team, and implemented supporting workflow automation across enterprise services.

Stack: LLM assistants, natural-language querying over platform data, backend integration, automation

Backend Engineer — Softylines

Mar 2021 – Feb 2024

Customer Engagement Platform

3 yrs

- Architected an **event-driven system with Kafka** (pub/sub, consumer groups, retry/DLQ) to decouple domains and support scalable asynchronous processing.
- Led a **10-engineer team** migrating a monolith to microservices, driving the plan from architecture through incremental rollout.
- Stabilised production by resolving high-impact incidents across messaging, webhooks, and real-time WebSocket surfaces, reducing recurring failure modes.
- Built and maintained Go and Node.js services balancing low-latency internal APIs against long-lived integrations and document-heavy workflows; owned CI/CD and mentored on API design and test strategy.

Stack: Go, Node.js, Kafka, WebSockets, PostgreSQL/MongoDB, Docker, observability

Selected Open Source

- **go-load-balancer** (Go, 12★) — production-ready HTTP load balancer.
- **go-cache** (Go, 6★) — high-performance Redis-like in-memory key-value store: transactions, lists, sets, JSON, pub/sub, TTL.
- **rag-gateway-service** (Python) — FastAPI gateway using OpenAI embeddings for semantic topic routing across RAG backends.
- **Pathfinding-Visualizer** (TypeScript, 7★) — interactive pathfinding algorithm visualiser.

Education

B.Sc. in Computer Science

June 2025

Horizon School of Digital Technologies